


# A Concurrent Hash Table for CPython

MSc Thesis Dissertation

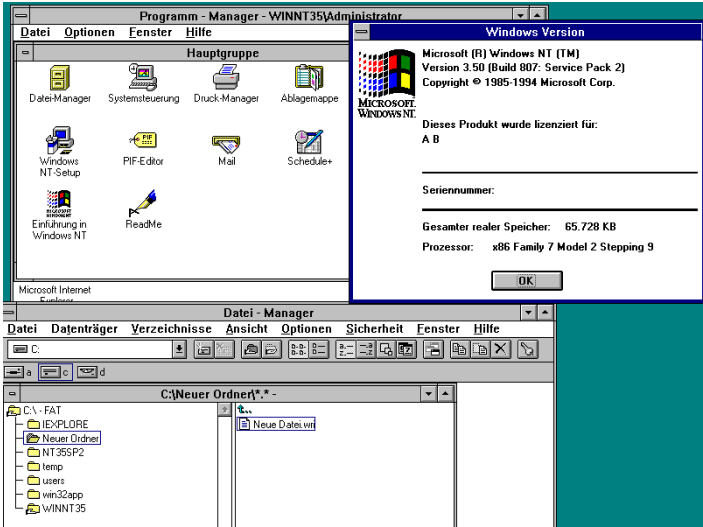
*Daniele Parmeggiani*

Università degli Studi di Trento, DISI

2024-10-16











## 3.13

### Free-threading

- \* Semaphore ✓
- \* CyclicBarrier ✓
- \* Lock ✓
- \* Concurrent**HashMap**
- \* ConcurrentLinked**Queue**
- \* CopyOnWriteArray**List**
- \* Atomic**Boolean**
- \* Atomic**Integer**
- \* Atomic**Reference**
- \* ...

Thread synchronization utilities for Python.

The *Cereus Greggii* is a flower native to Arizona, New Mexico, Texas, and some parts of northern Mexico.

This flower blooms just one summer night every year and in any given area, all these flowers bloom in synchrony.



<https://github.com/dpdani/cereggii>



Exposes a so-called *natural parallelism*.

High-level design:

- \* linear probing
- \* resizable
- \* double-hashing
- \* split index and data tables
- \* lock-free\*

Mostly inspired by:

- \* T. Maier's Grown concurrent hash table, and
- \* CPython's built-in sequential hash table (dict).




Figure: Grotz's migration process.<sup>1</sup>

<sup>1</sup>T. Maier et al., Concurrent Hash Tables: Fast and General(?), Fig. 1.

- \* Grown uses most-significant bits as position
- \* dict uses least-significant bits as position

- \* Grown uses most-significant bits as position
- \* dict uses least-significant bits as position
- \* AtomicDict:

- \* Growt uses most-significant bits as position
- \* dict uses least-significant bits as position
- \* AtomicDict:
  1. get an object's hash (as generated by CPython)

- \* Growt uses most-significant bits as position
- \* dict uses least-significant bits as position
- \* AtomicDict:
  1. get an object's hash (as generated by CPython)
  2. re-hash with CRC32

- \* Growt uses most-significant bits as position
- \* dict uses least-significant bits as position
- \* AtomicDict:
  1. get an object's hash (as generated by CPython)
  2. re-hash with CRC32
  3. use most-significant bits as position

- \* put keys and values into a separate *data* table
- \* actual hash table is an *index* over the data table



- \* put keys and values into a separate *data* table
- \* actual hash table is an *index* over the data table
  - \* size of an index slot: 1 to 8 bytes
    - \* dependent on capacity
  - \* size of a data entry: 32 bytes

- \* put keys and values into a separate *data* table
- \* actual hash table is an *index* over the data table
- \* so that:
  - \* the index stays sparse to better handle collisions
  - \* the large entry size decreases false sharing
- \* size of an index slot: 1 to 8 bytes
  - \* dependent on capacity
- \* size of a data entry: 32 bytes

## Definition

At least one thread can always make progress. (No locks  $\Rightarrow$  no deadlocks.)

## Definition

At least one thread can always make progress. (No locks  $\Rightarrow$  no deadlocks.)

- \* For AtomicDict:
  - \* no deadlocks!
  - \* (unless a resizing is in progress)
  - \* lock-free resizing is not practical
  - \* if you correctly set `min_size`, no resizing ever happens

## \* Partitioned iterations

```
d = AtomicDict(...)

def my_thread(thread_id, threads_count):
    for key, value in d.fast_iter(partitions=threads_count,
                                  this_partition=thread_id):
        do_some_stuff(key, value)
```

## \* Reduce

```
d = AtomicDict(...)

data = [
    ("red", 1),
    ("green", 42),
    ("blue", 3),
    ("red", 5),
]

def count(key, current, new):
    if current is NOT_FOUND:
        return new
    return current + new

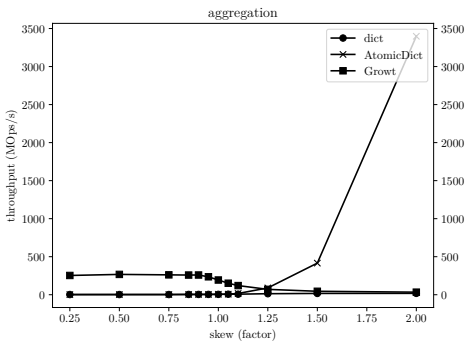
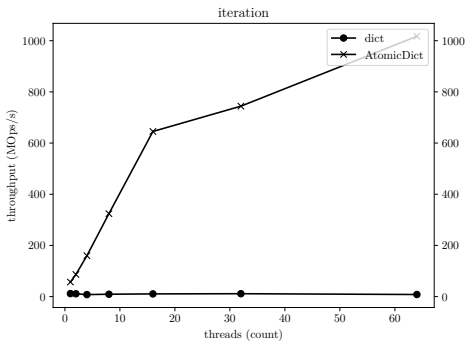
d.reduce(data, count)
```

## \* Batch Lookup

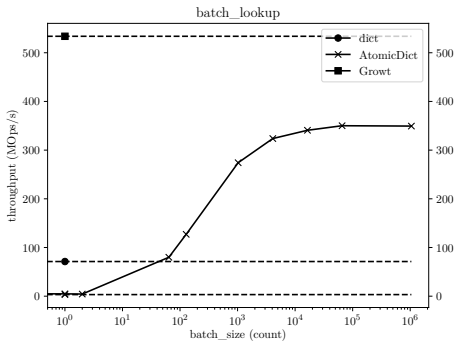
```
foo = AtomicDict({'a': 1, 'b': 2, 'c': 3})


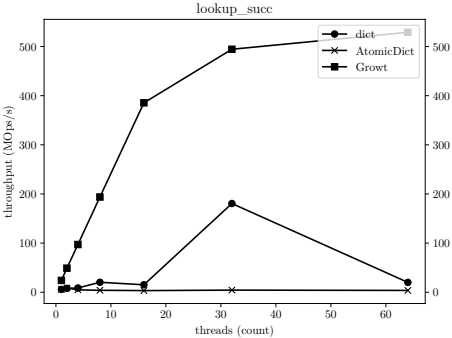
result = foo.batch_getitem({
    'a': None,
    'b': None,
    'f': None,
})

assert result == {
    'a': 1,
    'b': 2,
    'f': cereggii.NOT_FOUND,
}
```













Pittsburgh (USA), May 2024.

1. Python = language / CPython = interpreter
2. CPython and its GIL
  - \* historical rationale
3. Free-threading in CPython 3.13
  - \* per-object locks
4. Java has atomic data structures, Python doesn't
5. ceregii
  - \* atomic data structures for Python (currently: AtomicRef, AtomicInt, AtomicDict)

## 6. AtomicDict

- \* natural parallelism
- \* inspired by:
  - \* Maier's Growt (concurrent)
  - \* CPython's dict (sequential)
- \* linear probing
- \* resizable
- \* double-hashing
- \* split index and data tables
- \* lock-free\*
- \* partitioned iterations
- \* reduce
- \* batch lookup

## 7. Python Language Summit